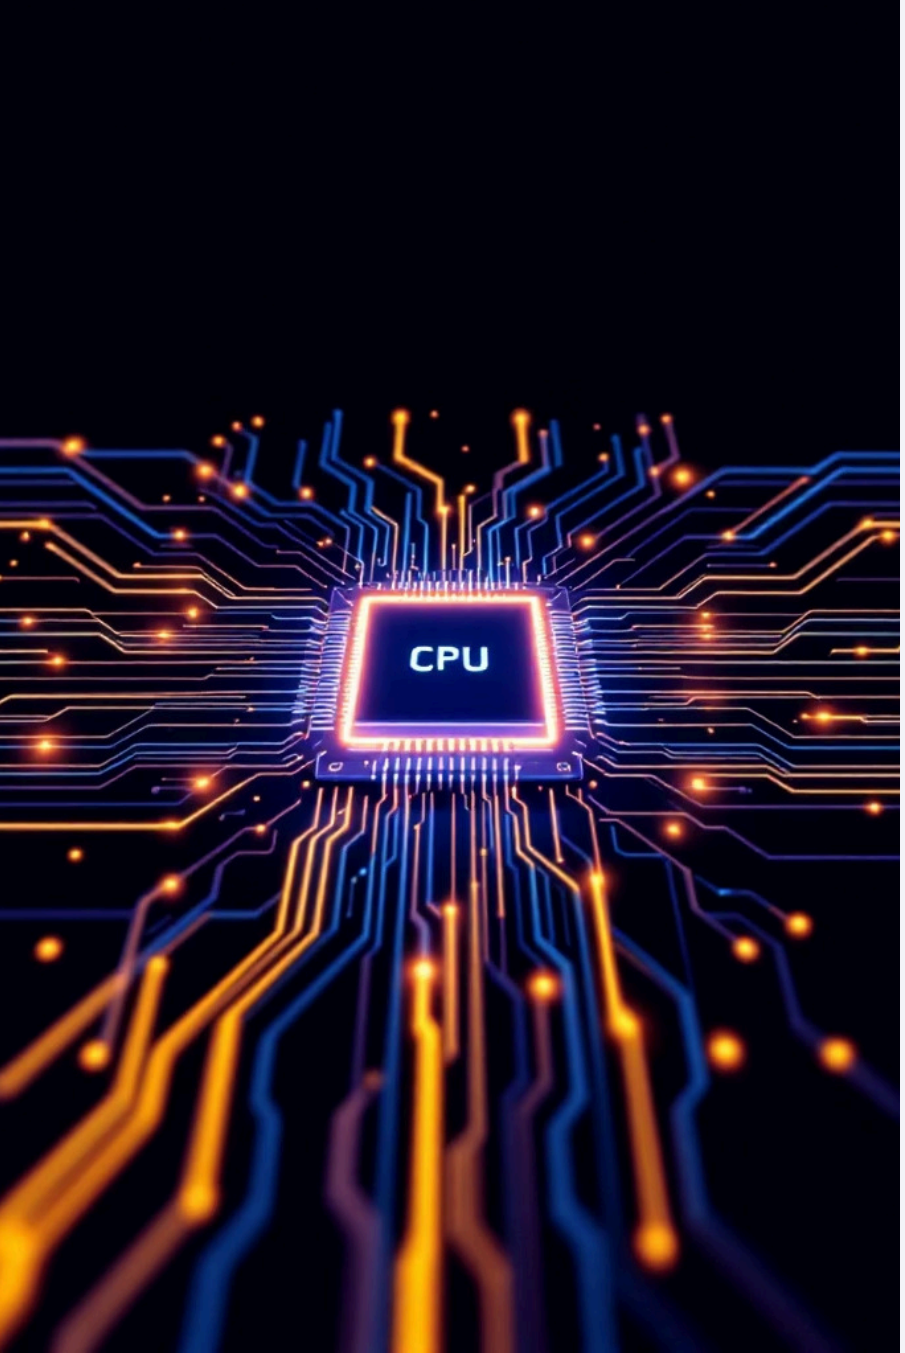

Understanding Computer Architecture

The Foundation of Modern Computing

From the silicon beneath your fingertips to the quantum frontiers ahead, computer architecture shapes how we compute, create, and connect.

Made with GAMMA

The Blueprint of Computing

What is Computer Architecture?

The science and art of designing the interface between hardware and software. It's the **blueprint** that defines how computers execute instructions and process data.

Why It Matters

Understanding architecture helps you write faster, cache-efficient code. It enables debugging complex concurrency issues and making informed hardware choices like ARM vs. x86 or CPU vs. GPU.

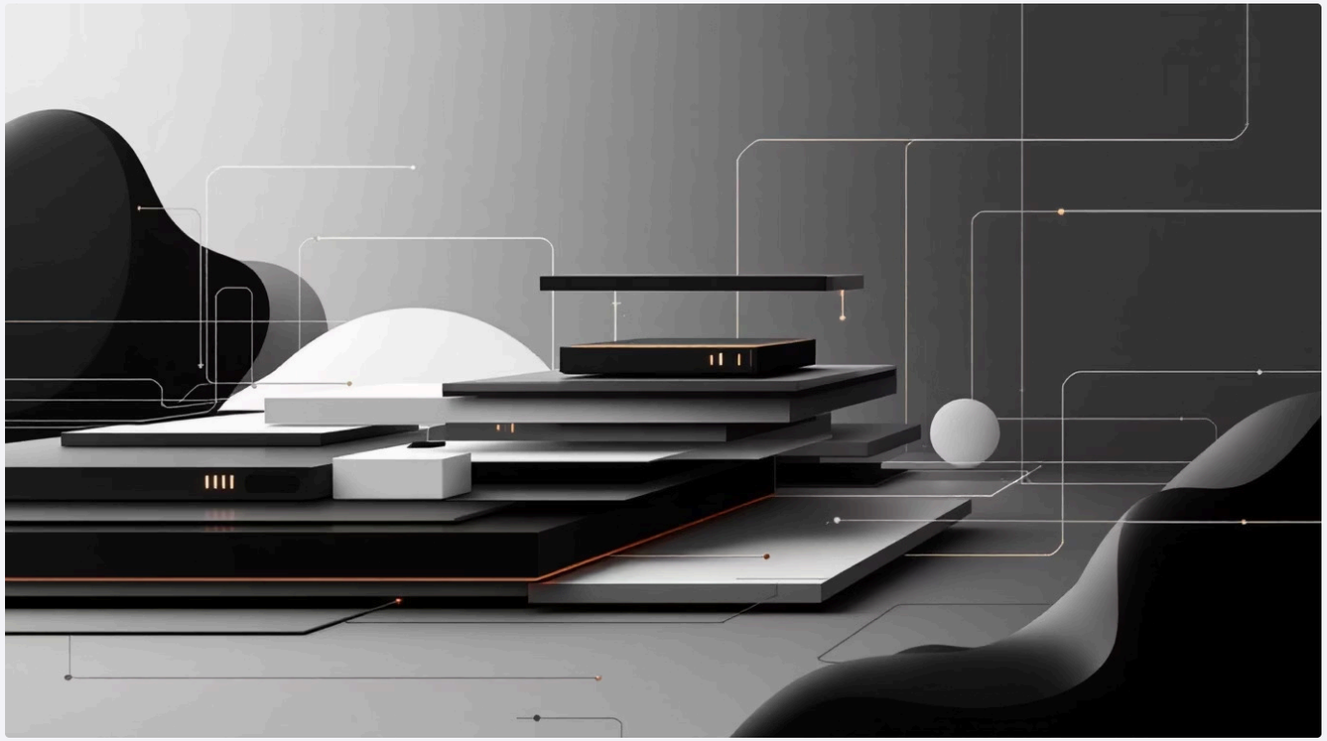

Instruction Set Architecture (ISA)

The programmer's view — the contract between software and hardware

Microarchitecture

How the ISA is actually implemented in silicon

Hardware

Physical logic gates, circuits, and chips

Made with GAMMA

RISC vs. CISC — Two Philosophies

Two competing approaches to instruction set design shaped modern computing, each with distinct strengths and trade-offs.

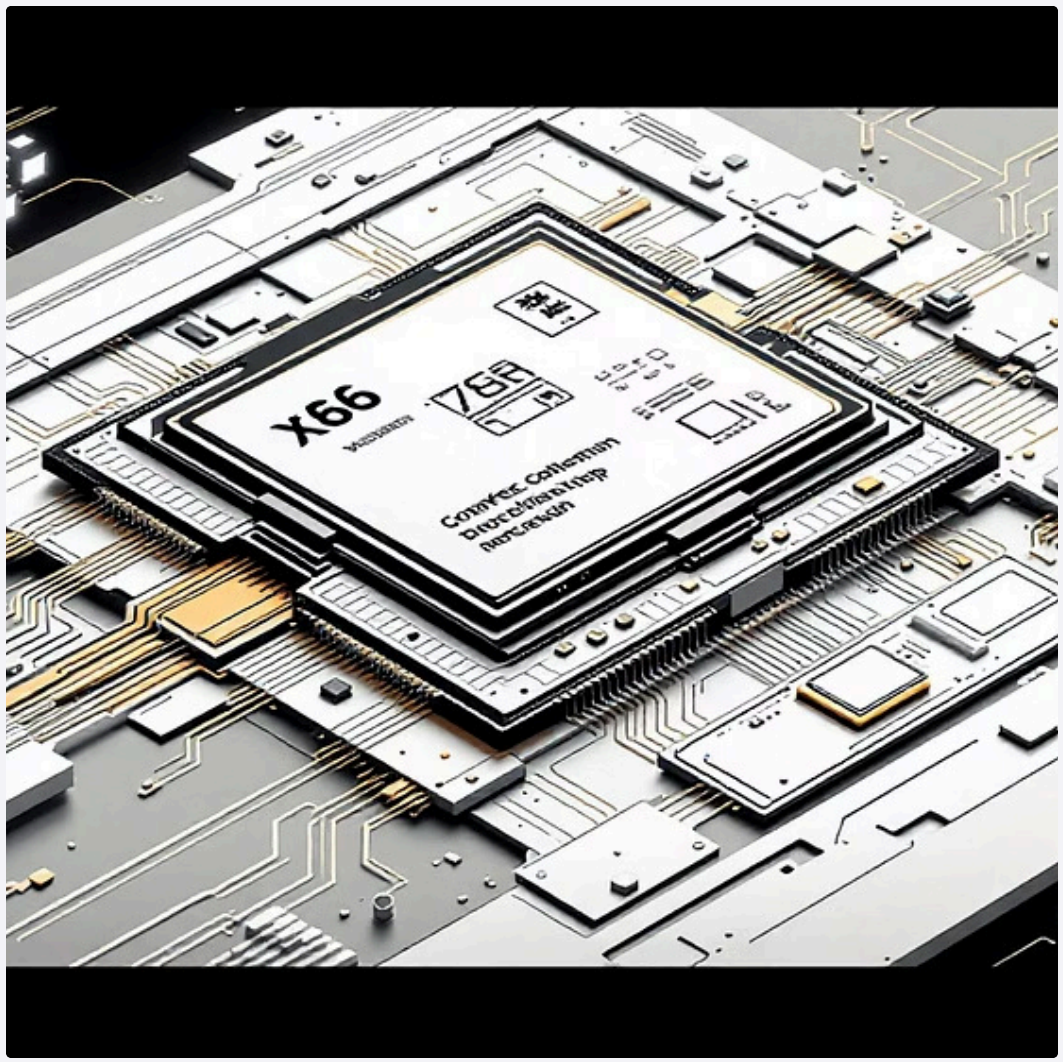

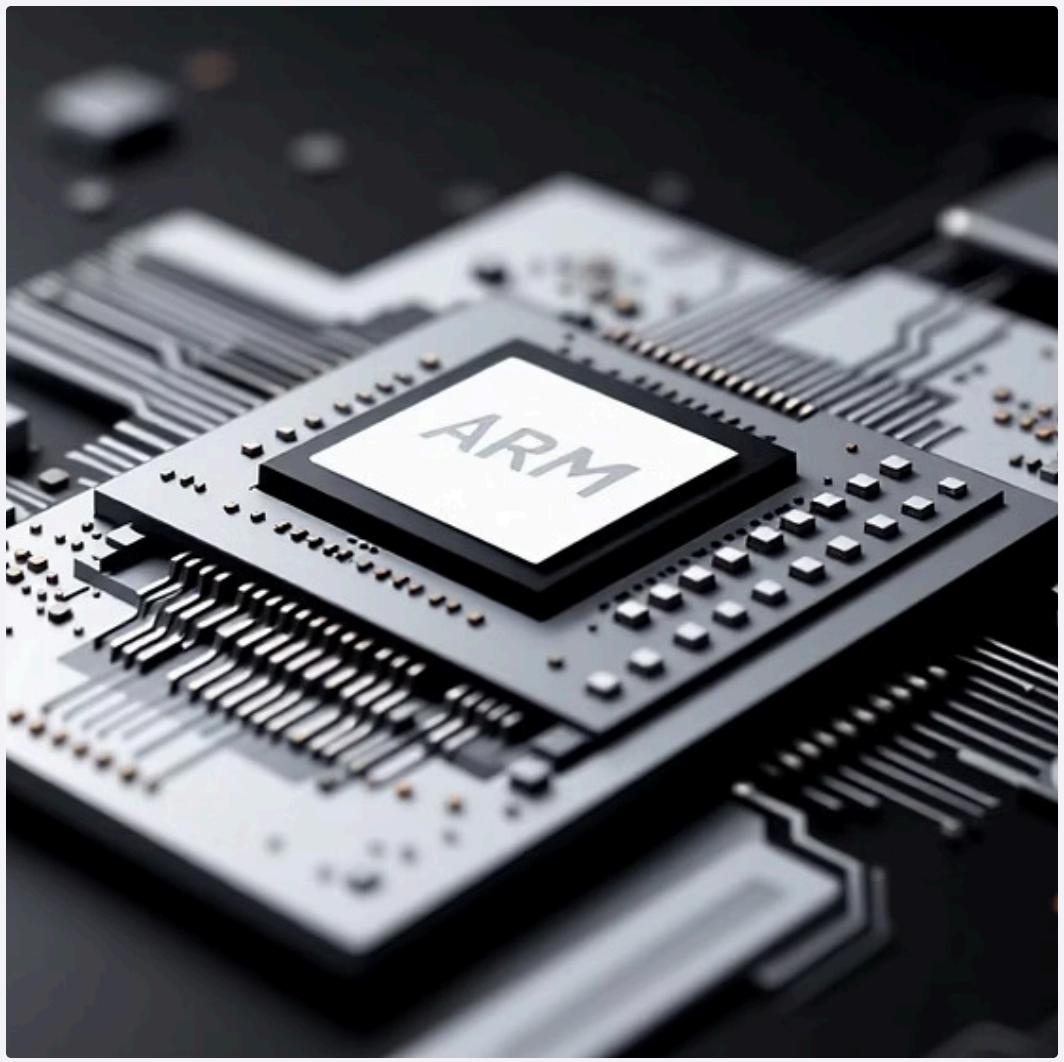

CISC (Complex Instruction Set Computer)

Examples: x86 (Intel, AMD)

- Complex, multi-step instructions
- Denser code, less memory usage
- Complex hardware, higher power

RISC (Reduced Instruction Set Computer)

Examples: ARM, RISC-V

- Simple, single-cycle instructions
- Simpler hardware, lower power
- Easier to pipeline and optimize

Modern Reality: CISC processors like x86 now internally decode complex instructions into RISC-like micro-operations — the lines are blurring.

Made with GAMMA

Doing More at Once

Pipelining: The Assembly Line Approach

Like an automotive assembly line, pipelining overlaps instruction execution stages. While one instruction executes, the next is decoded, and another is fetched — dramatically improving throughput.

Superscalar Execution

Multiple pipelines running in parallel, executing **more than one instruction per clock cycle**. This is the foundation of modern CPU performance.

Instruction-Level Parallelism (ILP)

The CPU's ability to identify and execute multiple independent instructions simultaneously, extracting maximum parallelism from sequential code.

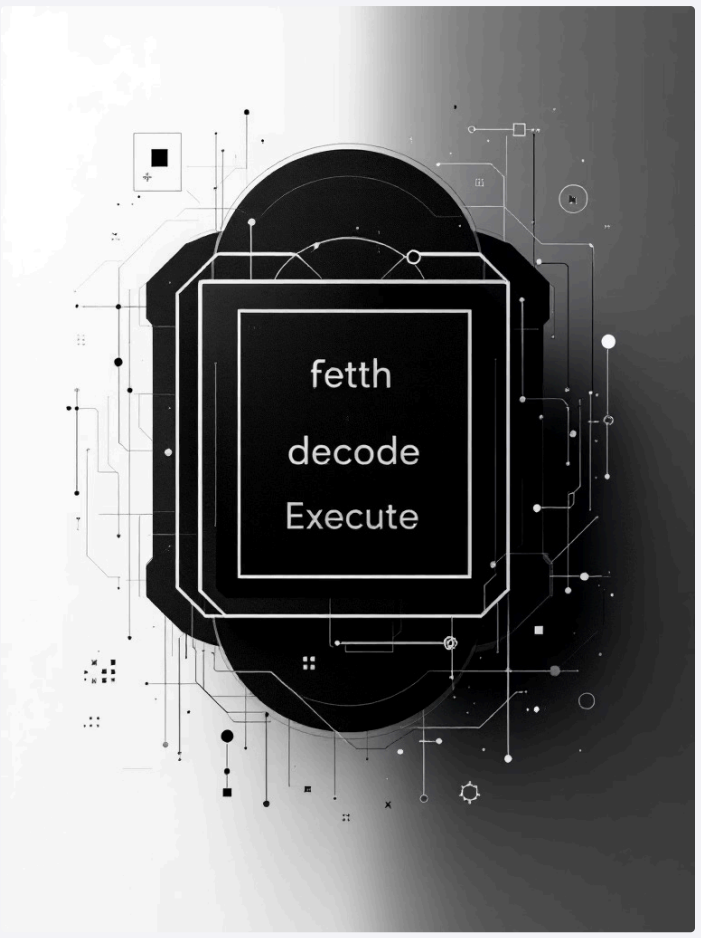

1. **Fetch**
   Retrieve instruction from memory
2. **Decode**
   Interpret instruction meaning
3. **Execute**
   Perform operation in ALU
4. **Memory**
   Load/store data to memory
5. **Writeback**
   Update register with result

Made with GAMMA

Inside the Central Processing Unit

Arithmetic Logic Unit (ALU)

Performs arithmetic operations (add, subtract) and logical operations (AND, OR, XOR). The computational engine of the CPU.

Control Unit

Directs data flow and instruction execution. Fetches instructions, decodes them, and orchestrates the ALU and other components.

Registers

The fastest memory — holds immediate data being processed. Limited in number but accessed in a single cycle.

Cache (L1/L2/L3)

Speed tiers bridging CPU and RAM. L1 is fastest/smallest, L3 is slower/larger. Reduces memory access latency dramatically.

Branch Prediction & Speculative Execution

Modern CPUs guess which path a branch will take to keep the pipeline full. This enables high performance but introduced security vulnerabilities like Spectre and Meltdown.

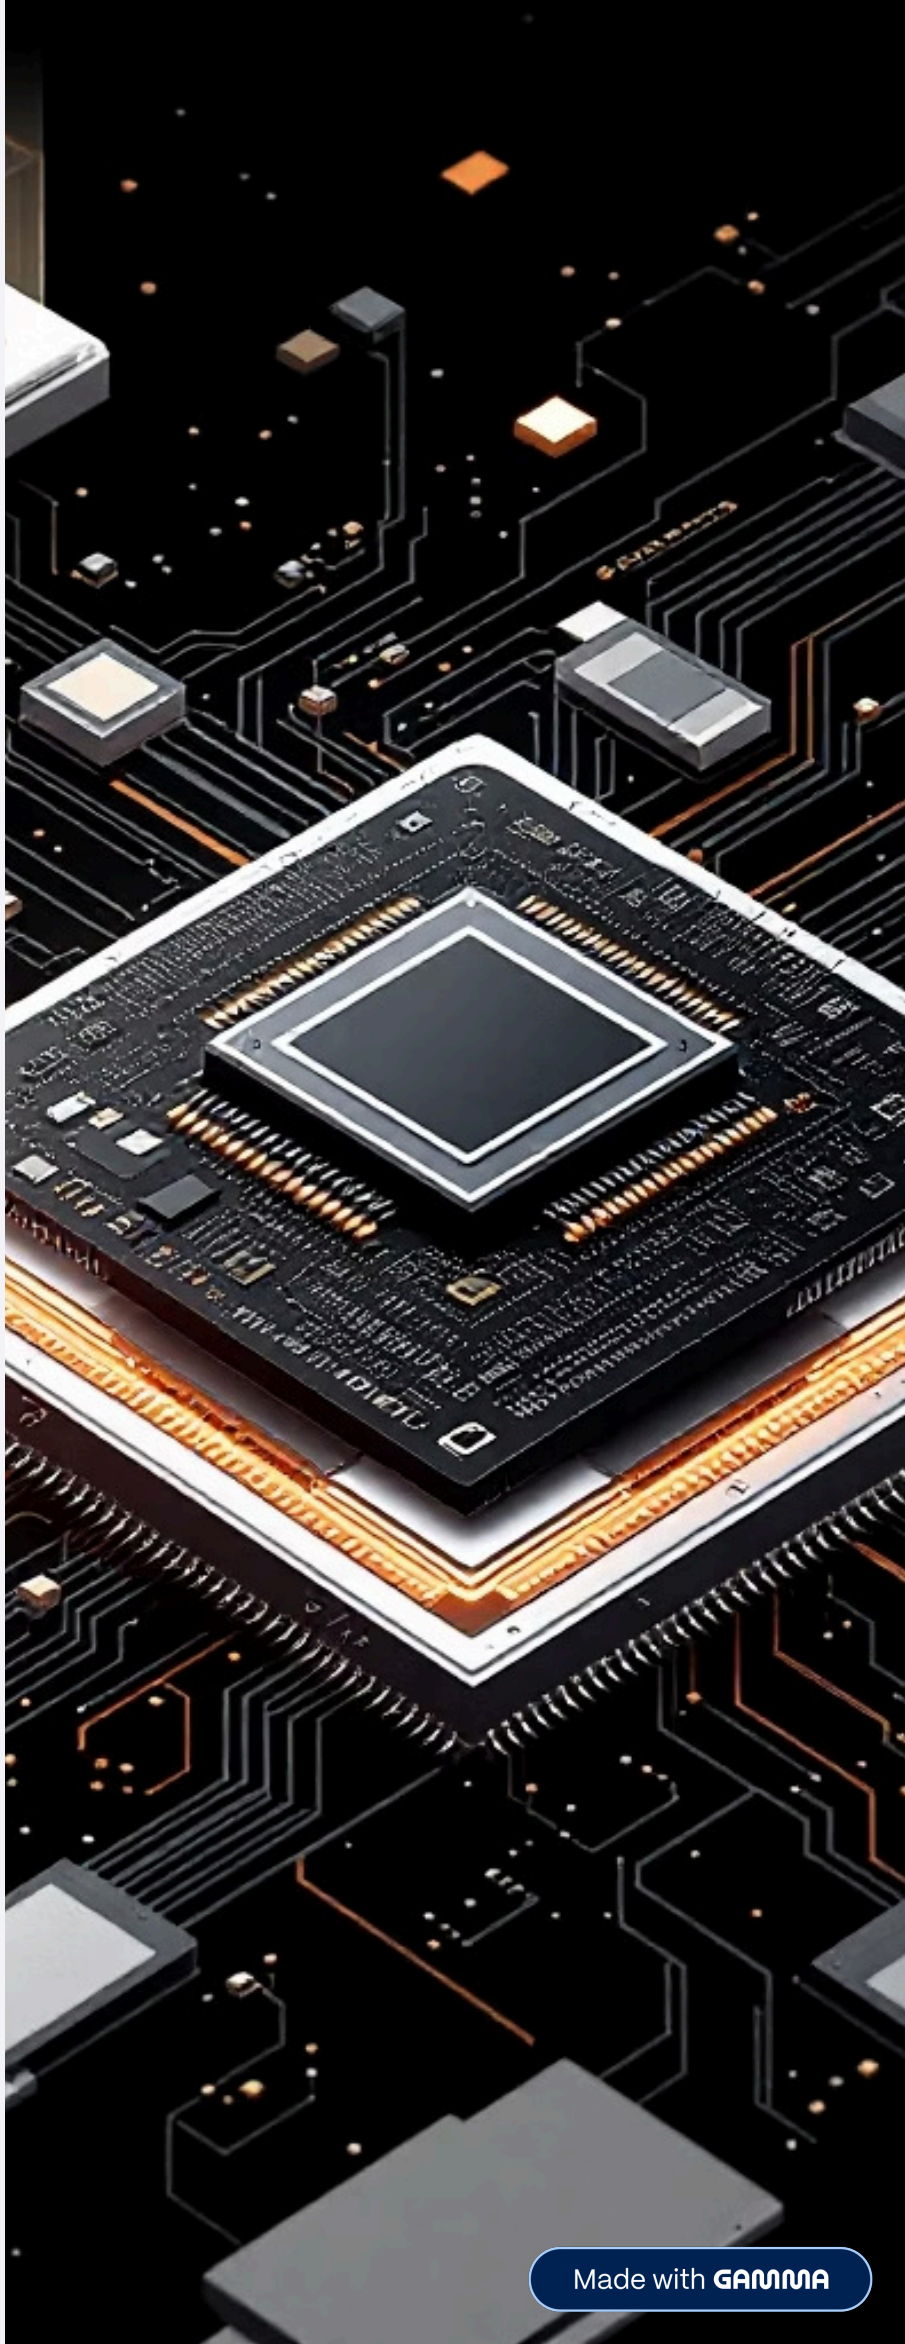

Made with GAMMA

Bridging the Speed Gap

Memory Hierarchy

Each level trades speed for capacity. Registers are fastest but smallest; cache bridges the gap; RAM provides working memory; storage offers massive capacity.

Virtual Memory

Abstracts physical RAM into a virtual address space for each process, providing isolation, the illusion of infinite memory, and efficient physical RAM use.

Paging

Memory divided into fixed-size pages, allowing swapping to disk when RAM is full. Enables systems to run programs larger than physical memory.

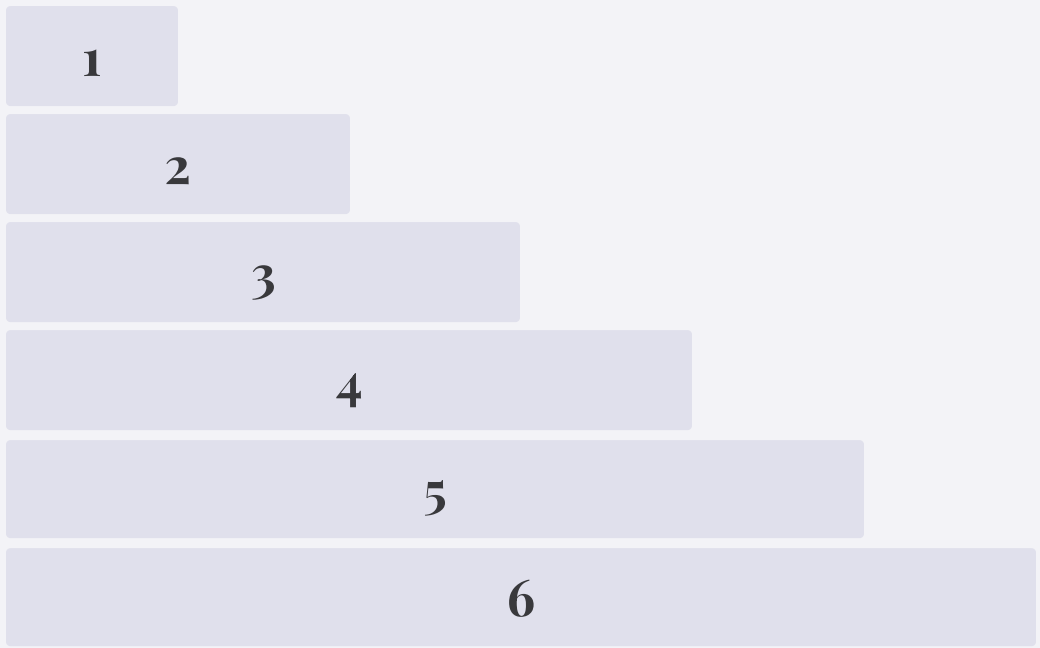

1 **Registers**

Fastest, smallest

2 **L1 Cache**

Ultra-fast

3 **L2 Cache**

Very fast

4 **L3 Cache**

Fast

5 **RAM**

Working memory

6 **SSD/HDD**

Massive storage

Made with GAMMA

Measuring Speed & Managing Heat

3.5GHz

Clock Speed

Cycles per second — important but not the full story

4.2

IPC

Instructions Per Cycle — work done per cycle

12

Core Count

Modern CPUs pack multiple cores for parallel processing

Performance Benchmarks

Standardized tests like **SPEC** and **Geekbench** provide fair comparisons across different architectures.

Power Wall

Higher clock speeds create exponentially more heat. The industry hit a wall around 2005, shifting to multi-core designs instead.

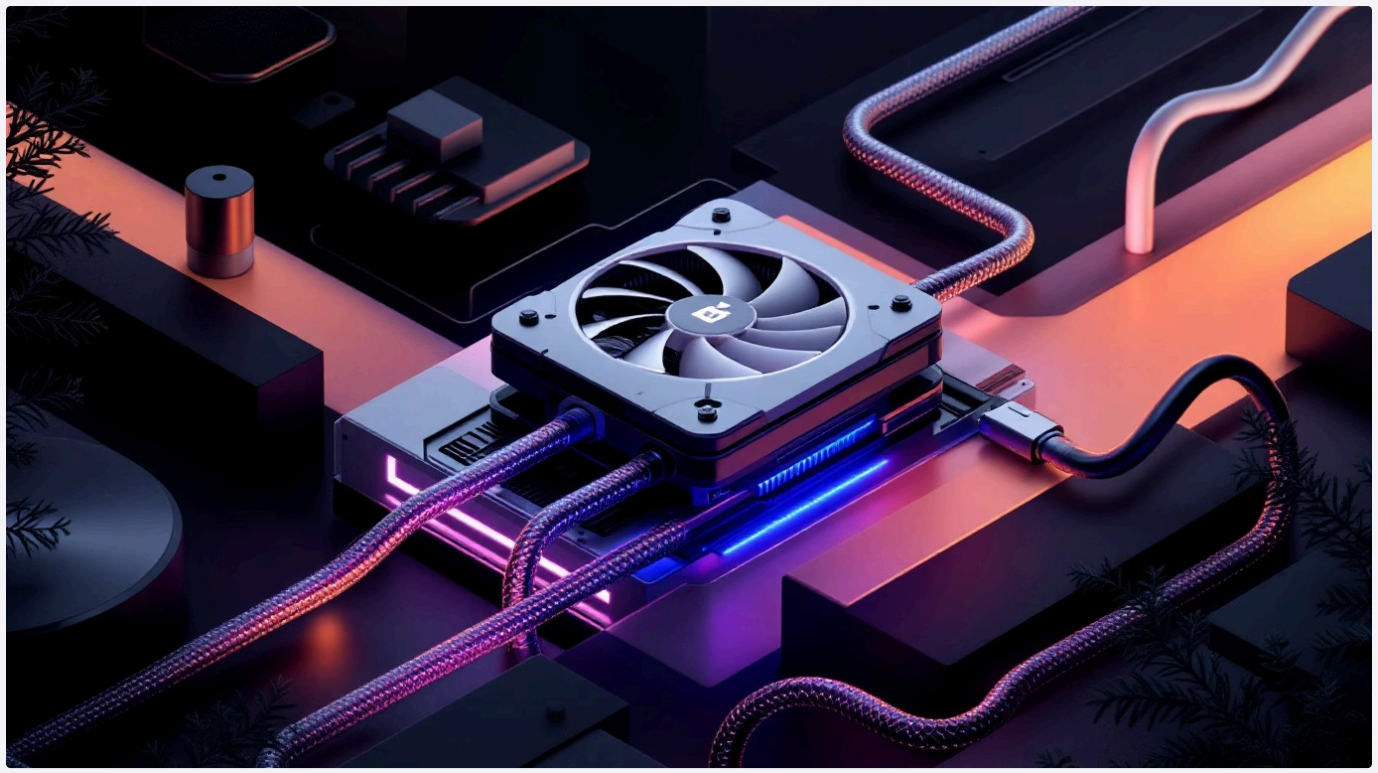

Cooling Solutions

Air cooling, liquid cooling, and even immersion cooling for extreme performance.

Dark Silicon: Parts of the chip must remain idle to avoid overheating — a fundamental limitation of modern processor design.

Made with GAMMA

Beyond Classical Architecture

Quantum Computing

Uses **qubits** with superposition and entanglement instead of classical bits. Excels at optimization, cryptography (Shor's algorithm), and quantum simulation.

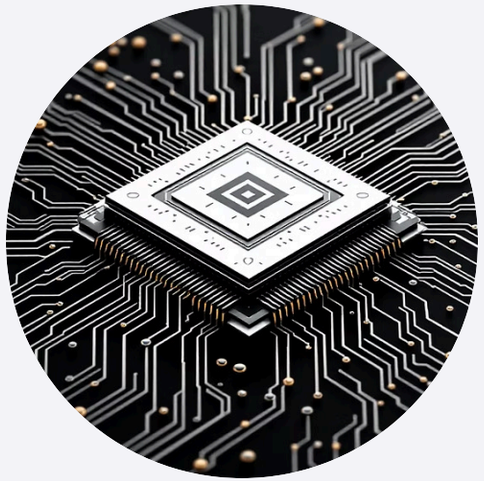

Neuromorphic Computing

Chips modeled after the human brain using spiking neural networks. Promises ultra-efficient AI processing.

Optical Computing

Using light instead of electricity for faster data transfer and lower power consumption.

The Evolution Continues

Computer architecture is evolving from faster clocks to specialized, parallel, and post-classical paradigms. **3D stacking** vertically integrates chiplets for reduced latency and higher density.

Classical CPUs will coexist with quantum processors for hybrid workloads — each optimized for different problem types.

Made with GAMMA